

Beyond Tokens: A Zero-Trust AuthZ Protocol for the Agentic Era

Nicola Gallo¹

Abstract—This paper addresses some of the security challenges posed by AI agents, framing them as part of the broader problem of distributed systems. It draws parallels with enterprise patterns for distributed systems and aims to lay the foundation for a new authorization (AuthZ) protocol. The key insight is that tokens, while useful, are not sufficient: authorization must move beyond a token-centric model. Instead, what is needed is a new zero-trust-authz protocol that, when combined with tokens, can support the development of next-generation AuthZ systems.

In practice, companies will continue to rely on standards such as OIDC and SAML, and OAuth will certainly remain an initiator for many flows. However, the real gap lies in what happens after token delivery — a space where existing protocols have paid little attention, but which is critical for the future of secure and dynamic authorization. The proposed model also leaves room for the integration of decentralized approaches, such as UCAN, ZCAP-LD, and other capability-based frameworks, enabling hybrid ecosystems that combine traditional standards with emerging decentralized authorization models.

I. INTRODUCTION

THE Agentic Era is emerging in industry discussions, as many recognize that certain assumptions no longer hold and past trade-offs are becoming increasingly fragile. Change can often be unsettling, especially when existing systems and business models are built on established approaches. Yet it may be time to start with a blank page and re-examine the compromises we have inherited.

The community should take this as a challenge to rethink the foundations of authorization, especially in the context of distributed systems and AI agents.

In the near future, for sure companies will continue to rely on OAuth to authenticate humans, and on solutions such as WIMSE or SPIFFE to authenticate non-human workloads. Yet at some point, we must stop patching legacy models and acknowledge that new architectures demand new foundations.

This new foundation can open entirely new markets. It can still leverage existing technologies, but with a fresh approach that enables innovation, new value creation, and the development of novel products. Industry has been built around identities, and most importantly this new foundation could enable the real adoption of AI agents within the enterprise.

The authentication layer is relatively solid, whether using centralized or decentralized approaches, but authorization still works only under certain trade-offs. These trade-offs cannot support AI agents and, more generally, distributed systems.

History of science teaches us that progress often comes from looking across disciplines, borrowing solutions, and applying them in new contexts, that is what this paper aims to do looking at Enterprise Patterns for distributed systems. It is also important to consider this in the context of Zero Trust Network Access

(ZTNA) and, more broadly, the Zero Trust model, which will be crucial for the future of secure access.

This paper adopts such an approach, challenging the core trade-off that the industry has relied on: the identity and HTTP-centric model. An AI agent protocol would need to include many other features; however, this paper focuses specifically on authorization (authz).

II. TOKENS IN THE AUTHORIZATION PERSPECTIVE

Tokens have worked very well for authentication (AuthN). An authentication token simply states: “*I am Nicola*”. If this statement is true now, it will remain true ever from now.

However, there are essentially two problems:

- 1) **Trust**: anyone can claim such a statement. To address this, tokens must be signed, giving rise to mechanisms such as JWTs.
- 2) **Replay**: anyone can replay the same token. The only option to mitigate this risk is to add an expiration time.

These are trade-offs and represent inherent limits of security. In practice, tokens exploit a PACELC-style trade-off: in the absence of partitions, they take advantage of the latency, consistency trade-off.

The PACELC theorem [9] states that in case of network partitioning (P) in a distributed computer system, one has to choose between availability (A) and consistency (C) (as per the CAP theorem), but else (E), even when the system is running normally in the absence of partitions, one has to choose between latency (L) and loss of consistency (C).

This trade-off becomes dangerous when tokens represent authorization (AuthZ). For example, “Nicola is entitled to a certain role now” may be true at one moment but false only a few seconds later. Here tokens become fragile: while they have worked well for authentication, it is a mistake to assume they must also be the golden standard for authorization.

Below we present a more formal representation of this problem.

A. Temporal Validity and Risk

Let T_a be the validity period of an authentication token and T_z the validity of an authorization. To ensure security:

$$T_z \ll T_a, \quad \text{ideally } T_z \rightarrow 0$$

In practice, T_z should be kept as short as feasible, though operational trade-offs may apply.

The risk of compromise grows with validity:

$$R(T) = f(T), \quad f'(T) > 0$$

Thus, long-lived static authorization tokens increase exposure.

¹Nicola Gallo, Software Architect at Nitro Agility. This work was carried out as part of professional activities within Nitro Agility and should be understood in that professional context. nicola.gallo@nitroagility.com

B. Transitive Delegation

A delegation between A and B is denoted:

$$D_i : A \rightarrow B$$

A chain of n delegations is:

$$D_1 \circ D_2 \circ \dots \circ D_n$$

If a token explicitly encodes a single linear chain of n delegations, its size grows linearly:

$$O(n)$$

However, when delegation supports branching, conditions, or attenuation, the number of possible valid paths increases combinatorially. Encoding all possible transitive delegations into a static token leads to exponential growth in the worst case:

$$O(2^n)$$

This is not scalable. Therefore, authorization must be reconstructed dynamically at request time.

C. Context and Runtime Validity

System state evolves with time:

$$S(t) = \{E, W, P\}$$

where E are entities, W workloads, and P policies.

At request time t_r , validity requires:

$$\exists p \in P \text{ valid for } (e, w) \mid t = t_r$$

Static tokens cannot guarantee consistency with the live system state $S(t_r)$.

D. Trust Model

Authorization is a function of the live context:

$$T(E, W, P, t) = \begin{cases} 1 & \text{if authorized at time } t \\ 0 & \text{otherwise} \end{cases}$$

Static tokens implicitly encode authorization state at issuance time t_0 , and reuse assumes that the same decision holds for all later times:

$$T(E, W, P, t_0) = T(E, W, P, t), \forall t > t_0$$

which is false in dynamic systems.

E. Conclusion

Token-based authorization is:

- 1) **Insecure**: $R(T)$ increases with T .
- 2) **Non-scalable**: delegation graphs may induce combinatorial explosion (up to $O(2^n)$).
- 3) **Inconsistent**: $S(t) \neq S(t_0)$ in dynamic systems.

Therefore, **authorization must be computed on-demand**, consistent with Zero Trust principles.

F. Practical Considerations

The model above is a theoretical ideal; in practice it is almost impossible to achieve perfectly. As in many areas of security, trade-offs must be made to balance feasibility and protection. What matters is not the invention of a “better token” but the design of a more sophisticated authorization engine capable of assessing multiple sources of trust — including tokens, capabilities (e.g. zcaps), contextual signals, and other dynamic inputs — in a best-effort approach.

The key point is that the next generation of authorization protocols cannot remain token-centric. Instead, authorization must emerge from a richer machinery that integrates diverse identities and trust sources, and can even operate *without* tokens altogether.

III. THE INDUSTRY EVOLUTION TO ZERO TRUST

CLOUD computing has reshaped how companies build software, shifting the focus to scalable solutions that leverage elastic infrastructure. Applications are increasingly containerized and managed with platforms such as Kubernetes, while serverless computing has also gained traction.

Being inherently distributed, cloud-native systems face the challenges outlined by the CAP theorem: architects must balance Partition Tolerance, Consistency, and Availability. This has led to designs that relax strong consistency in favor of availability through Eventual Consistency.

The CAP theorem [1] states that, in the presence of a network partition, a system must trade off between Consistency and Availability. Eventual Consistency aligns with PA (Partition Tolerance and Availability), allowing a system to remain responsive during partitions while deferring full consistency until recovery.




Fig. 1. CAP Theorem

Another major trend is the adoption of asynchronous patterns and event-streaming frameworks such as Apache Kafka, which handle continuous flows of events without a clear beginning or end. In this setting, orchestration has proven complex and difficult to scale, whereas choreography-based systems often provide better scalability.

Figure 2 shows a typical cloud architecture: clients authenticate with a centralized Identity Provider, which issues a token subsequently used to access server resources. Two widely adopted protocols here are OpenID [4] for authentication and OAuth [5] for authorization.

Within these protocols, JSON Web Tokens (JWT) [6] have become central. Compact and digitally signed, they ensure mes-

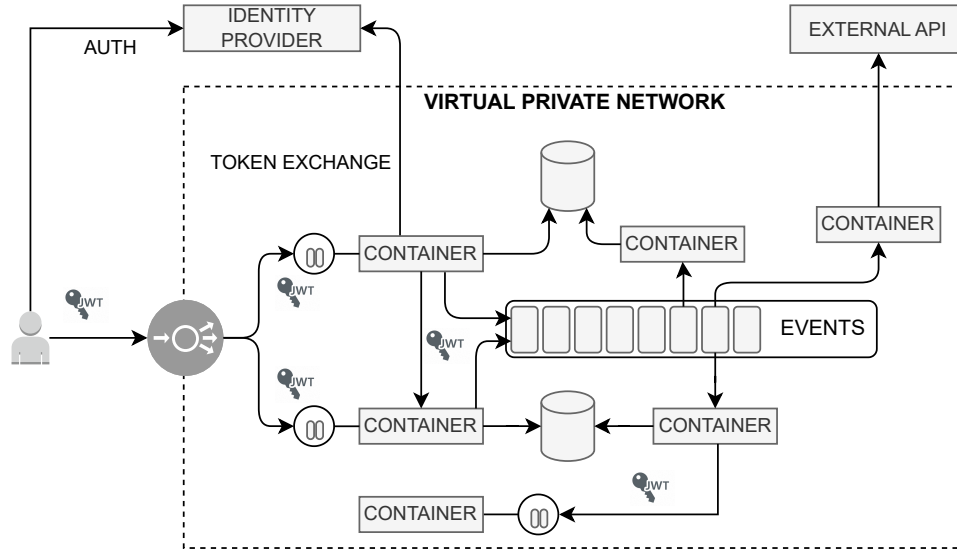


Fig. 2. Cloud Architecture

sage integrity and enable secure communication between client and server.

Access control remains a critical concern. Models such as Role-Based Access Control (RBAC) and Attribute-Based Access Control (ABAC) have paved the way for Policy-Based Access Control (PBAC) [7].

A further challenge arises in transport-layer communication, for example with Kafka and stream-processing systems, where it is not always feasible to propagate a token.

without clear perimeters. In cloud-native, multi-tenant, and hybrid environments, workloads, services, and agents continuously interact across organizational and network boundaries. In such scenarios, the very assumption of a fixed perimeter for validation no longer holds, making token-based authorization even more fragile.

It follows that the current model is not suitable for AI agents, which require dynamic, context-aware, and continuously validated authorization. The goal is not to replace existing protocols but to complement them, addressing the gaps they have historically overlooked.

ACKNOWLEDGMENTS

This text was originally written by the author in English. As a non-native English speaker, AI-based writing tools were employed solely to refine grammar, clarity, and fluency in order to speed up the writing process and reduce language errors. All ideas, arguments, and original contributions are entirely the author's own.

The views and proposals expressed in this paper are intended for discussion and research purposes only. They should not be interpreted as finalized standards or production-ready specifications.

REFERENCES

- [1] Brewer, Eric. (2000). Towards robust distributed systems. PODC. 7. 10.1145/343477.343502.
- [2] V. Velepucha and P. Flores, "A Survey on Microservices Architecture: Principles, Patterns and Migration Challenges," in IEEE Access, vol. 11, pp. 88339-88358, 2023, doi: 10.1109/ACCESS.2023.3305687.
- [3] V. Vernon, Domain-Driven Design Distilled. Reading, MA, USA: Addison-Wesley Professional, 2016.
- [4] [OpenID.2.0] OpenID Foundation, "OpenID Authentication 2.0," December 2007.
- [5] [RFC6749] M. Jones, D. Hardt, "The OAuth 2.0 Authorization Framework," October 2012.
- [6] [RFC6750] M. Jones, D. Hardt, "The OAuth 2.0 Authorization Framework: Bearer Token Usage", October 2012
- [7] L. Zhi, W. Jing, C. Xiao-su and J. Lian-xing, "Research on Policy-based Access Control Model," 2009 International Conference on Networks Security, Wireless Communications and Trusted Computing, Wuhan, China, 2009, pp. 164-167, doi: 10.1109/NSWCTC.2009.313.




Fig. 3. Building inside a Perimeter

This exposes a key limitation of the current model: it works at the perimeter, where tokens are validated at delivery, but what happens inside the enterprise afterward is not standardized. Authorization is essentially assumed to remain valid once accepted at the boundary, and the system continues to treat it as true for the rest of the time. While this trade-off was once acceptable, it cannot be sustained in environments that must support AI agents.

Moreover, modern distributed systems increasingly operate

- [8] Nicola Gallo, Antonio Radesca, "A Multi-Account and Multi-Tenant Policy-Based Access Control (PBAC) Approach for Distributed Systems Augmented with Risk Scores Generation".
- [9] D. J. Abadi, "Consistency Tradeoffs in Modern Distributed Database System Design: CAP is Only Part of the Story," *IEEE Computer*, vol. 45, no. 2, pp. 37–42, Feb. 2012. DOI: 10.1109/MC.2012.33